



University of Colorado **Boulder**

Department of Computer Science

CSCI 5622: Machine Learning

Chenhao Tan

Lecture 1: Introduction

Slides adapted from Jordan Boyd-Graber, Chris Ketelsen

Basic Information

- Course location and time: HUMN 1B80, 5:00-6:15pm (MW)
- Instructor: Chenhao Tan
- Teaching assistant: Han Liu
- Course manager: Kodur Krishna Chaitanya
- Graduate student staff: Brian Groenke
- Course website: <https://github.com/BoulderDS/CSCI5622-Machine-Learning>

About me

- Research interests: human-centered machine learning, natural language processing, computational social science
- <https://chenhaot.com>

Outline

- An overview of machine learning
- Syllabus
- Administtrivia

Machine learning is everywhere!

Reading comprehension

The Normans (Norman: Nourmands; French: Normands; Latin: Normanni) were the people who in the 10th and 11th centuries gave their name to Normandy, a region in France. [...]

Q: In what country is Normandy located?

A: France

Three lidar systems

A forward facing camera

Radar sensors

Self-driving sensors






6




7:05

1



● ALPHAGO
00:10:29



● LEE SEDOL
00:01:00

③



⑤



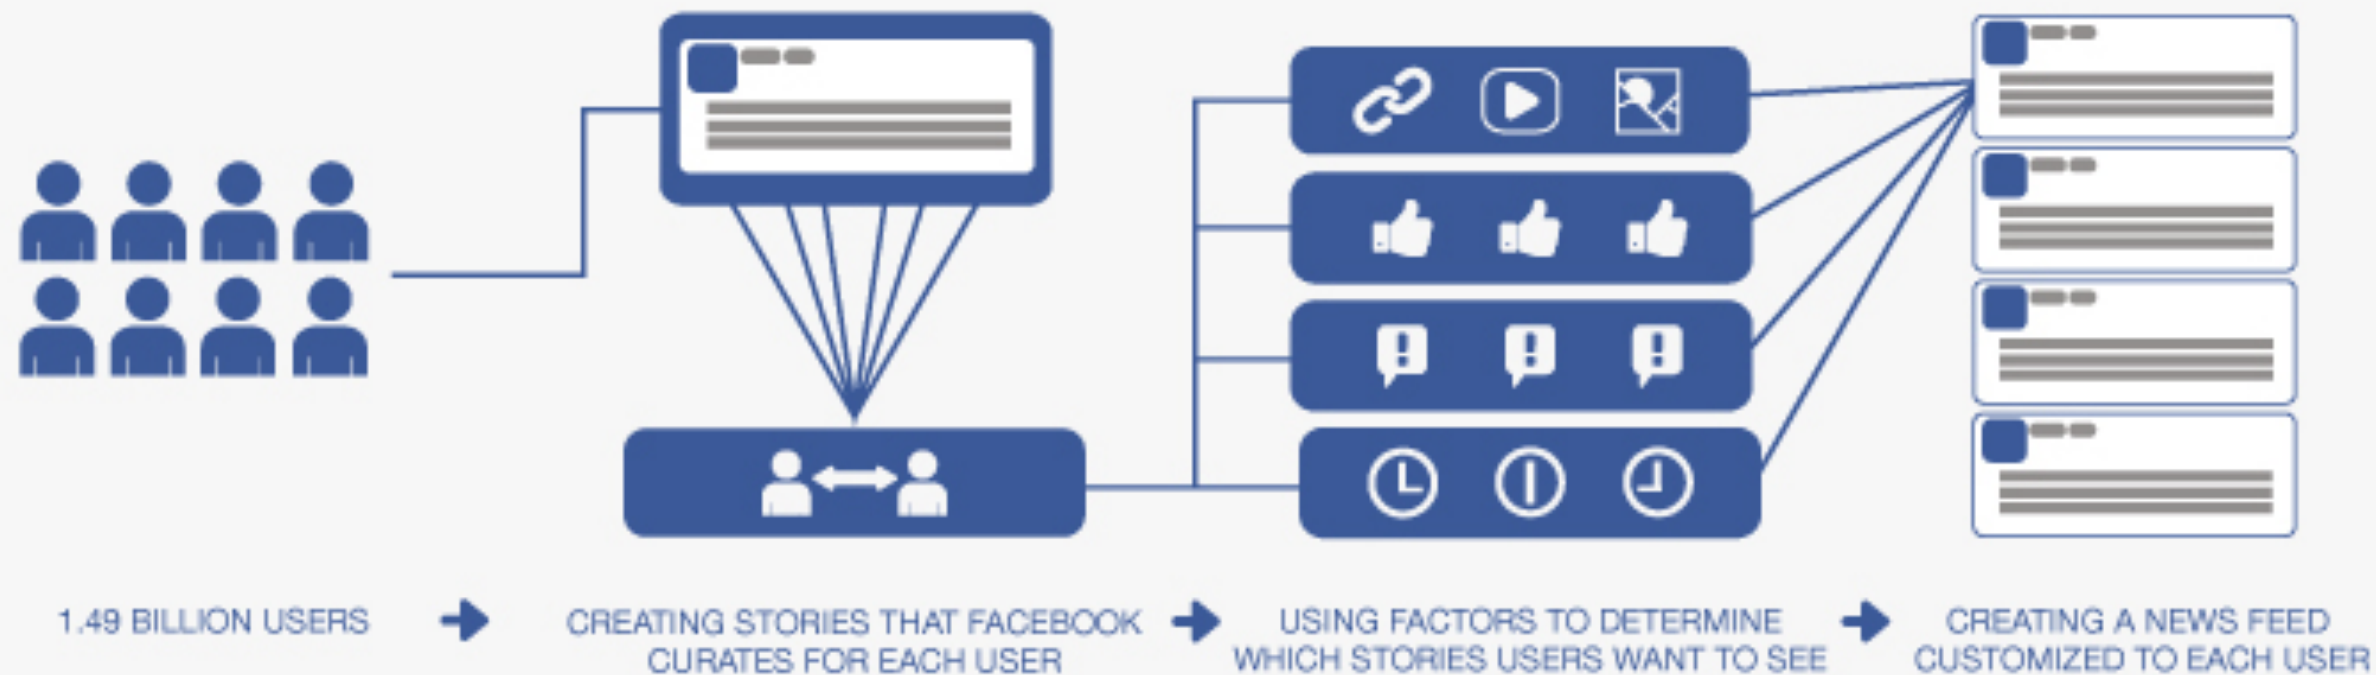
⑤



⑤



FACEBOOK'S NEWS FEED & USER EXPERIENCE



Which tweet will be retweeted more?



Food trucks are the epitome of small independently owned LOCAL businesses! Help keep them going! Sign the petition bit.ly/P6GYCq



I know at some point you've have been saved from hunger by our rolling food trucks friends. Let's help support them! bit.ly/P6GYCq

Which tweet will be retweeted more?



cactus_music
@cactus_music



Food trucks are the epitome of small independently owned LOCAL businesses! Help keep them going! Sign the petition bit.ly/P6GYCq



cactus_music
@cactus_music

I know at some point you've have been saved from hunger by our rolling food trucks friends. Let's help support them! bit.ly/P6GYCq

[GET CHART](#)[COMPARE](#)[EVENTS ▾](#)[TECHNICAL INDICATORS ▾](#)[CHART SETTINGS ▾](#)[RESET](#)

Aug 29, 2012 2:59 PM - 3:04 PM EDT: ■ ^DJA 4416.03



© 2012 Yahoo! Inc.

■ Volume: 571,700





ANNALS OF MEDICINE APRIL 3, 2017 ISSUE

A.I. VERSUS M.D.

What happens when diagnosis is automated?

By Siddhartha Mukherjee



Machine learning is everywhere!

- Smart city
- Entertainment
- Social
- Finance
- Medical

Email me to introduce yourself, one of your core values, and a machine learning application you care about.

What is machine learning?

One definition (Mitchell):


A computer program is said to learn from

- experience E with respect to some class of
- tasks T and
- performance measure P ,

If its performance at tasks in T , as measured by P , improves with experience E .

Let us apply this to classic tasks in machine learning!

ImageNet/Object recognition



- T : identifying objects in an image
- E : tons of images with annotated objects
- P : how often the objects are identified correctly

Sentiment analysis



Hutzler #571 Banana Slicer

The only banana slicer you will ever need.

Gourmac's easy-to-use Banana Slicer provides a quick solution to slice a banana uniformly each and every time. Simply press the slicer on a peeled banana and the work is done. Safe, fun and easy for children to use. Kids just love eating bananas with this as their favorite kitchen tool. The Banana Slicer may also be used as a quick way to add healthy bananas to breakfast cereal or to make uniform slices for a fruit salad or ice cream dessert.


★ ★ ☆ ☆ ☆ **Angle is wrong**

By Jim Anderson on August 1, 2012

I tried the banana slicer and found it unacceptable. As shown in the picture, the slicer is curved from left to right. All of my bananas are bent the other way.

- T : deciding whether a review is positive or negative
- E : reviews with ratings
- P : how often the sentiment is predicted correctly

Movie recommendation



- T : recommend movies
- E : movie watching history and movie ratings
- P : future ratings of users? User active time on website? User subscription periods?

Why do we want machines to learn?

- Intellectually satisfying!
- Helping us solve problems

Automate tasks that we know how to perform

- Object recognition
- Driving

Emulate human ability

Explore tasks that we don't know how to perform

- Predict message popularity
- Movie recommendation
- Newsfeed ranking

Extend human ability

Human-centered machine learning

Help humans solve critical societal challenges

- Medical diagnosis
- Recidivism prediction
- Personalized education

How does machine learning work as of today?

- Collect or happen upon data (X , experience in the previous definition)
- Analyze it to find patterns
- Use those patterns to perform some task (T)

We will study algorithms that find and exploit patterns in data.

We will study algorithms that find and exploit patterns in data.

- Goal: fluency in thinking about modern machine learning algorithms
- We will learn about
 - When to use them
 - The assumptions they make about data
 - Their capabilities and their limitations
 - Theoretical guarantees
- We will learn a language and process for solving data analysis problems.

Supervised vs. unsupervised methods



Hutzler #571 Banana Slicer

The only banana slicer you will ever need.

Gourmac's easy-to-use Banana Slicer provides a quick solution to slice a banana uniformly each and every time. Simply press the slicer on a peeled banana and the work is done. Safe, fun and easy for children to use. Kids just love eating bananas with this as their favorite kitchen tool. The Banana Slicer may also be used as a quick way to add healthy bananas to breakfast cereal or to make uniform slices for a fruit salad or ice cream dessert.

Data: X

Labels: Y

- **Supervised methods** find patterns in **fully observed** data and then try to predict something from **partially observed** data.
- For example, in sentiment analysis, after learning something from annotated reviews, we want to take new reviews and automatically identify sentiments.

Supervised vs. unsupervised methods



Data: X Hidden
Structure: Z

- **Unsupervised methods** find **hidden structure** in data, structure that we can never formally observe.
- For example, modeling topics from a collection of scientific papers; evaluation is usually more difficult.

Outline

- An overview of machine learning
- **Syllabus**
- Administrivia

Preliminary schedule

- Supervised learning
- Unsupervised learning
- Maybe learning theory
- Ethics
- <https://github.com/BoulderDS/CSCI5622-Machine-Learning/blob/master/info/schedule.md>

Prerequisites


- Programming language: Python
- Math background:
 - Probability
 - Linear algebra
 - Calculus
 - *Information theory*
 - *Optimization*

Prerequisites

- Programming language: Python
- Math background:
 - Probability
 - Linear algebra
 - Calculus
 - *Information theory*
 - *Optimization*

Highly recommend
trying the prerequisite
exam on Moodle!

Course textbook



Outline

- An overview of machine learning
- Syllabus
- **Administrivia**

Contact information

- Course on github:
<https://github.com/BoulderDS/CSCI5622-Machine-Learning/>
- Piazza (all Q&A should happen here):
<https://piazza.com/class/jza7ib732nu1yz>
- Moodle (all grading happens here):
<https://moodle.cs.colorado.edu/course/view.php?id=1231>

Grading policy

- Homework (35%)
- Exams (30%)
- Final project (30%)
- Participation (5%)

Grading policy

- Homework (35%) (Count the highest four, and NO LATE SUBMISSIONS.)
- Exams (30%)
- Final project (30%)
- Participation (5%)

Final project

- Competition or replication or research
 - Kaggle competitions
 - Replicate a research paper
 - Work on something related to your research

Final project

- Group formation due: Oct 4
- Final project proposal due: Oct 25
- Final project midpoint report (this can be a short update): Nov 15
- Hackathon: Dec 4
- Final project presentation: Dec 11
- Final project report due: Dec 13

Learning is a two-way process

- Private emails
 - Only for feedback, all Q&A should happen through Piazza
- Periodic survey

Q & A

- Email me to introduce yourself, one of your core values, and a machine learning application you care about.